


Scheduling Policies

- Batch: FIFO, SJF, STCF
- Proportional sharing: RR, Lottery, MLFQ
- MLFQ (demotion policy, promotion policy)
 - • Rule 1: If $\text{Priority}(A) > \text{Priority}(B)$, A runs (B doesn't). • Rule 2: If $\text{Priority}(A) = \text{Priority}(B)$, A & B run in RR. • Rule 3: When a job enters the system, it is placed at the highest priority (the topmost queue). • Rule 4a: If a job uses up an entire time slice while running, its priority is reduced (i.e., it moves down one queue). • Rule 4b: If a job gives up the CPU before the time slice is up, it stays at the same priority level., also Rule 5: Priority boost and gaming tolerance.

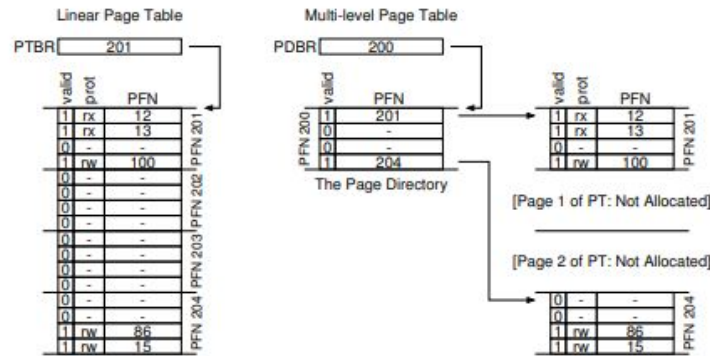


- I/O blocking == sleeping

Virtualizing Memory / MMU mechanisms/address translation schemes

- Base & Bounds (2 registers)
 - translate via addition
 - - B&B are two registers that are used to translate registers
 - - they provide an offset to the address when you're converting from virtual address to physical address - BASE
 - - Bounds - applied before or after (depends) that says what is the max address that you are allowed to use.
 - Privileged mode is needed so programs can't do bad things, also exceptions
- Segmentation ($N \times 2$ registers, typically $N=2,3,4$)
 - translate via choice+addition
 - Segmentation is like BnB but multiple, one for code, one for data ...
- Paging
 - fixed-sized pages (4K typical)
 - translate via lookup+offset
 - (virtual) page numbers vs (physical) frame numbers
 - uses a translation table, not mathematical relation
 - page table entries (PTE) have control bits:
 - valid
 - readable / writable / executable
 - (dirty, accessed, present, ...)
- Multi-level paging
 - page directory vs page tables
 - - virtual boxes and physical boxes of same size
 - - each virtual mem unit has an offset from start of box and box number
 - - Page table say has 5 entries in it, which is a pointer to one of the physical memory spaces.
 - - $T[1] = 6$, meaning VM box 1 = found in PM box 6
 - - there also would be flags n such.
 - - each entry is a word, the "6" is the really important stuff, and the rest of the bits can be used to store other information such as validity, user status, read/write, and other features.
 - - if we think about multiple level tables,

- - so instead we have an array of PTE (page table entries, but some PTEs refer to other PTE which then actually point to where they want to in physical mem.



-
- Translation Lookaside Buffers (TLBs)
 - address space IDs (ASIDs), they are stored on this so it is quicker and don't need to view page table every time. This is to "Make the Common Case Fast"
 - VPN | PFN | bits

Virtual Memory Problems

- Base & bounds requires contiguous physical memory allocation
- Segmentation requires helps trivial cases, but ignores: loadable code, threads, shared memory, ...
- B&B, Segmentation require "malloc-like" memory management
- Paging has excessive overhead due to page table size
- Paging doubles all memory accesses (thus significantly harms performance)
- Multi-level paging has even more memory access overhead (thus ...ugh!)
- TLBs significantly impact context switches and violate isolation

Multi-level Paging w/ TLB Features

- relatively low space and time overhead (but violates isolation, information can leak)
- shared code between processes (DLLs, SOs)
- demand loading of code
- shared memory between processes (low overhead inter-process communication)
- automatic stack growth
- null/stray pointer detection
- memory-mapped I/O (e.g., a file mapped into process address space)
- hardware emulation (e.g., old display hardware)
- virtual memory (see below)

Virtual memory (swapping to disk)

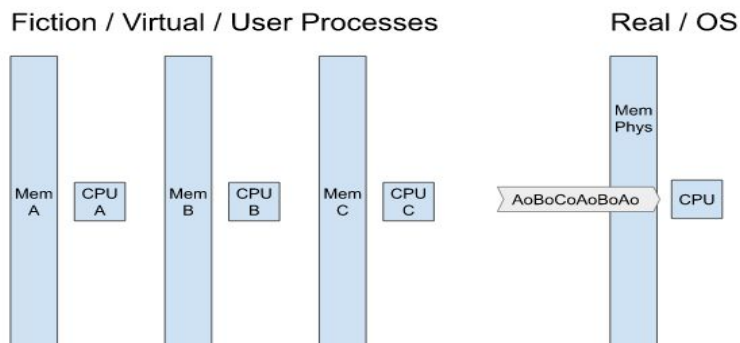
- swap space (relative access cost)
- page fault handling, present page control bit, blocking
- background page writer
- policies: optimal (prescient) FIFO, random, LRU, LFU, clock
- historical (stateful) accounting costs, page control bits: accessed, dirty
- terms: thrashing, working set

Concurrency Synchronization (30 31 32)

- Problems
 - Atomicity - interruption of operations assumed to be atomic
 - Ordering - order is different than expected
 - Deadlock - waiting without moving

- Livelock - waiting while looping
- User-Space Sync Mechanisms
 - None - Simple, atomic
 - Block Interrupts - Prevents scheduling during atomic instruction sequences, bad OS authority
 - Instructions - Used to implement others
 - Spin-wait -waste of CPU
- OS Sync Services
 - Thread join - Very high-level order management
 - Mutex - Informs the scheduler of critical sections
 - lock(a) - begin crit section : unlock(a)- end crit section
 - Cond Var - identifies threads waiting for a resource
 - signal(a) : wait(a)
 - Semaphore - Counts and queues up threads waiting for a resource (Producer/Consumer Model)
 - post(a) : wait(a)
 - Contributed by Dijkstra, allows data to be transferred between threads
- Common Sync Problems
 - Producer/consumer - making and using data, such as games
 - Read/Writer - simultaneously access data for reading, others exclusively write, databases
 - Dining Philosophers - threads require multiple resources simultaneously, order of obtaining matters
- Thread safe Data structures
 - Concurrent access across threads, not by OS, Usually via mutex/locks, ideally atomic sync
 - Push and pop could implement producer/consumer without visible sync code

Threads - Multiple CPUs per Process



- Thread Models
 - User-space threads(old unix) - OS schedules between processes, proces scheudles between threads, blocked thread blocks the whole process
 - Kernal Threads - Threads = Processes(Linux) same address space tho, Threads as unique OS Resource(windows) but kernal schedules threads and processes

xv6 per-process state

```

struct proc {
    uint sz;                // Size of process memory (bytes)
    pde_t* pgdir;           // Page table
    char *kstack;           // Bottom of kernel stack for this process
    enum procstate state;   // Process state
    int pid;                // Process ID
    struct proc *parent;     // Parent process
    struct trapframe *tf;    // Trap frame for current syscall
    struct context *context; // switch() here to run process
    void *chan;             // If non-zero, sleeping on chan
    int killed;             // If non-zero, have been killed ?????
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;       // Current directory
    char name[16];          // Process name (debugging)
};

```